

Optimizations

- Compilation for Embedded Processors -

Peter Marwedel
TU Dortmund
Informatik 12
Germany

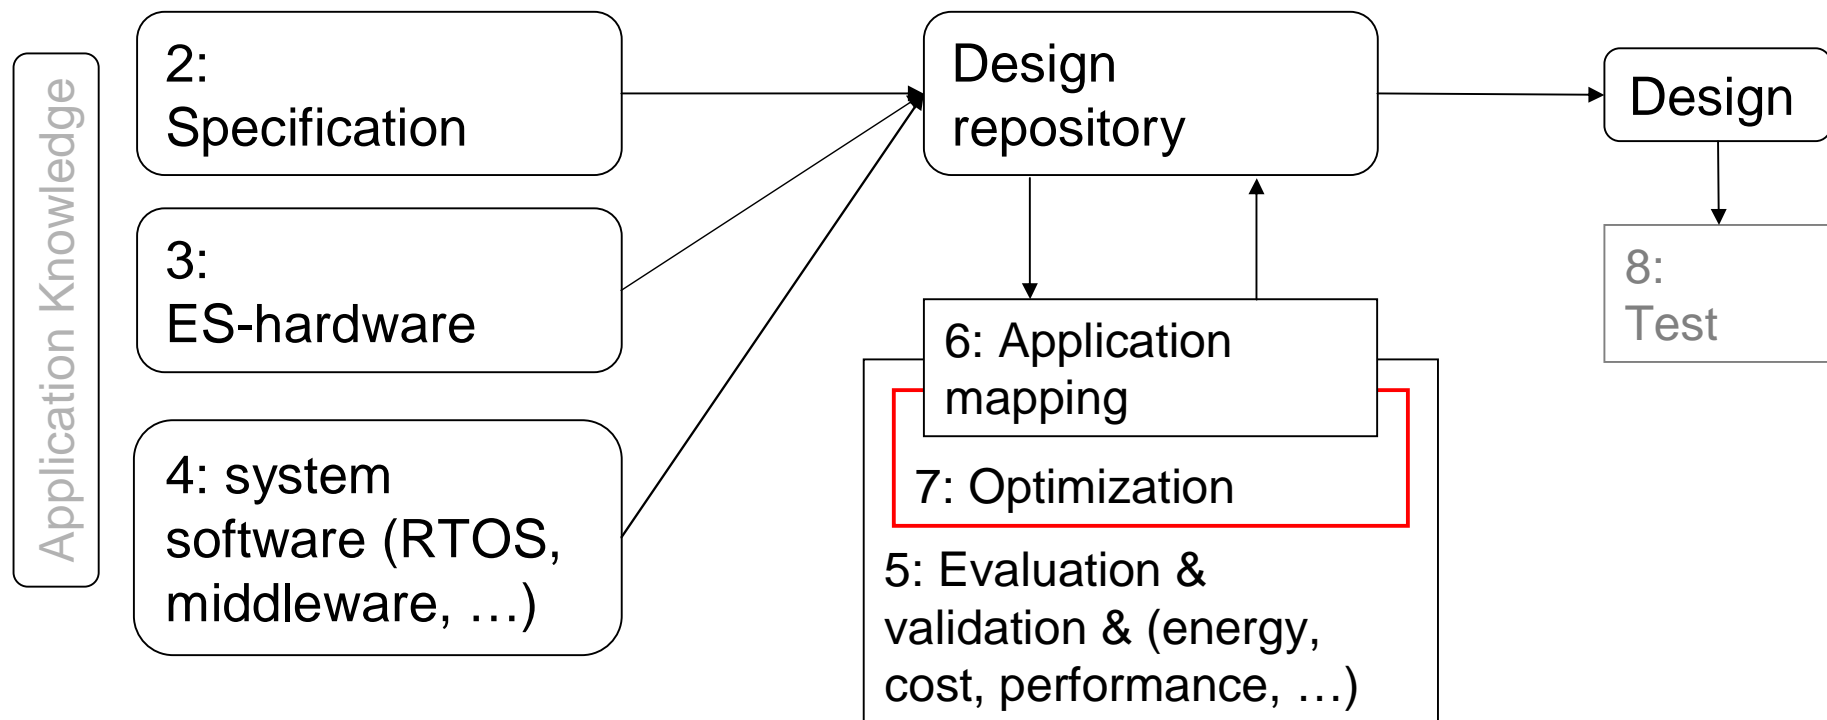
2013年 01 月 23 日



© Springer, 2010

These slides use Microsoft clip arts. Microsoft copyright restrictions apply.

Structure of this course



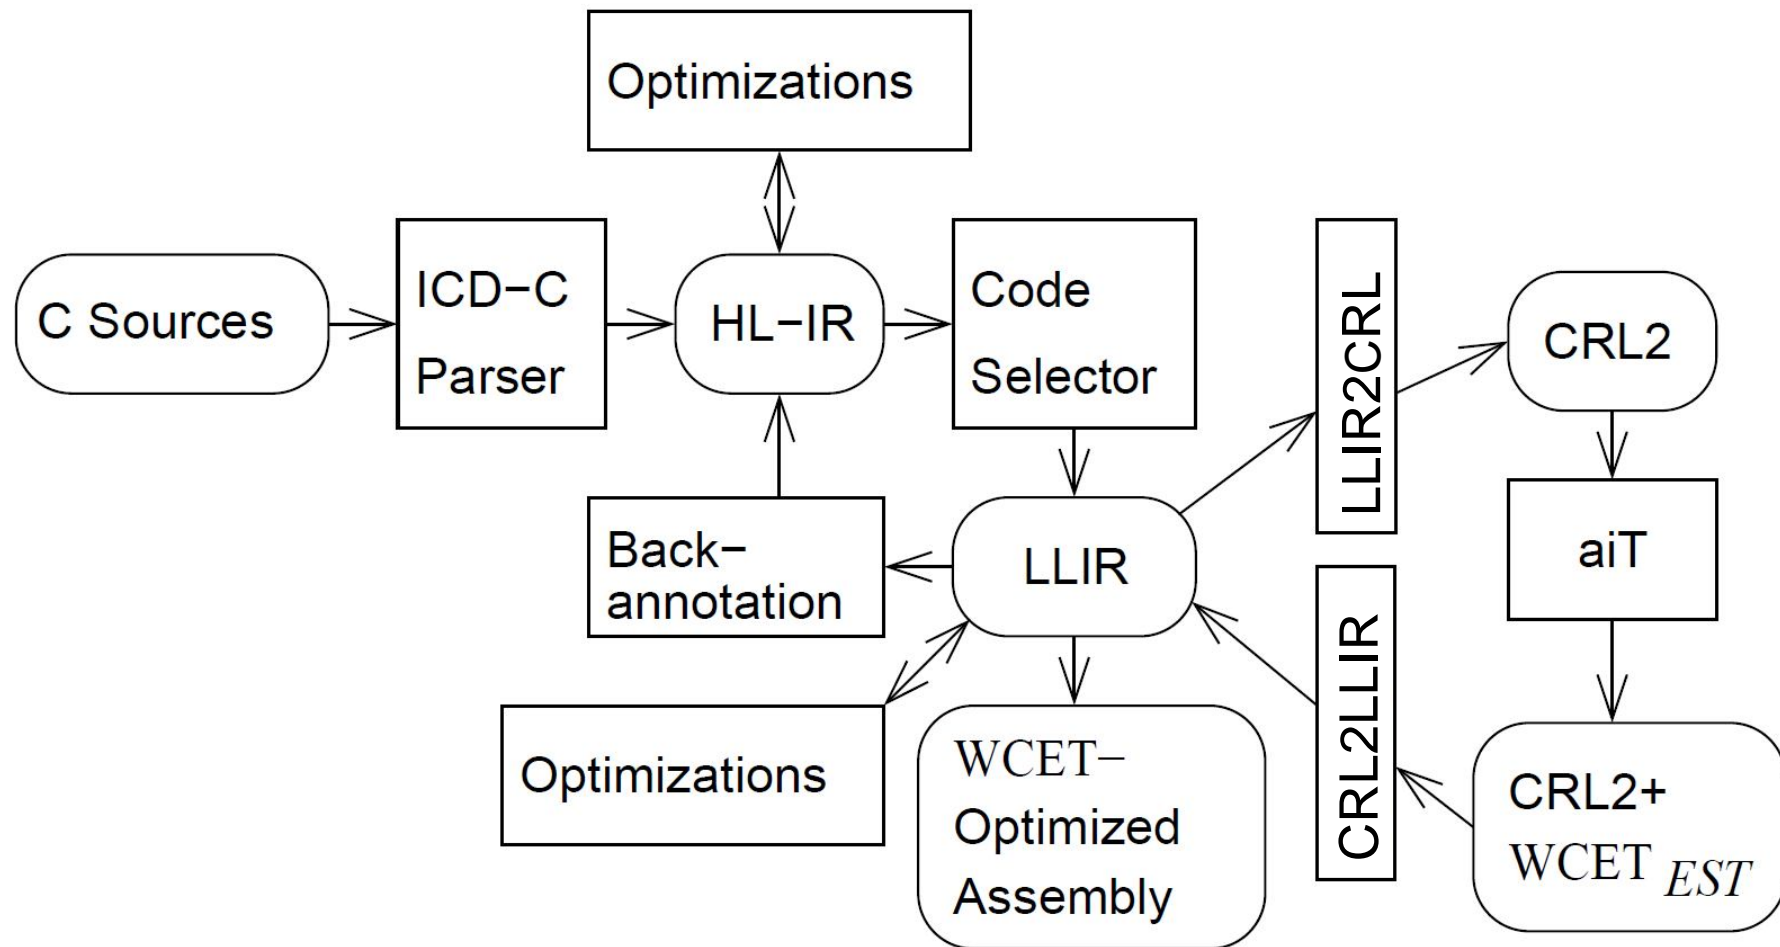
Numbers denote sequence of chapters

Reconciling compilers and timing analysis


Compilers mostly unaware of execution times

- Execution times are checked in a “trial-and-error” loop:
{try: compile – run – check – error: change}*
 - Timing analysis based on measurements is not safe
- ☞ Integration of safe, static timing analysis into compiler
- Getting rid of loops (if everything works well)
 - Safe timing verification
 - Potential for optimizing for the WCET

Structure of WCC (WCET-aware C-compiler)



Results for WCET-aware register allocation



The offset assignment problem

Peter Marwedel
TU Dortmund
Informatik 12
Germany

2010/01/13

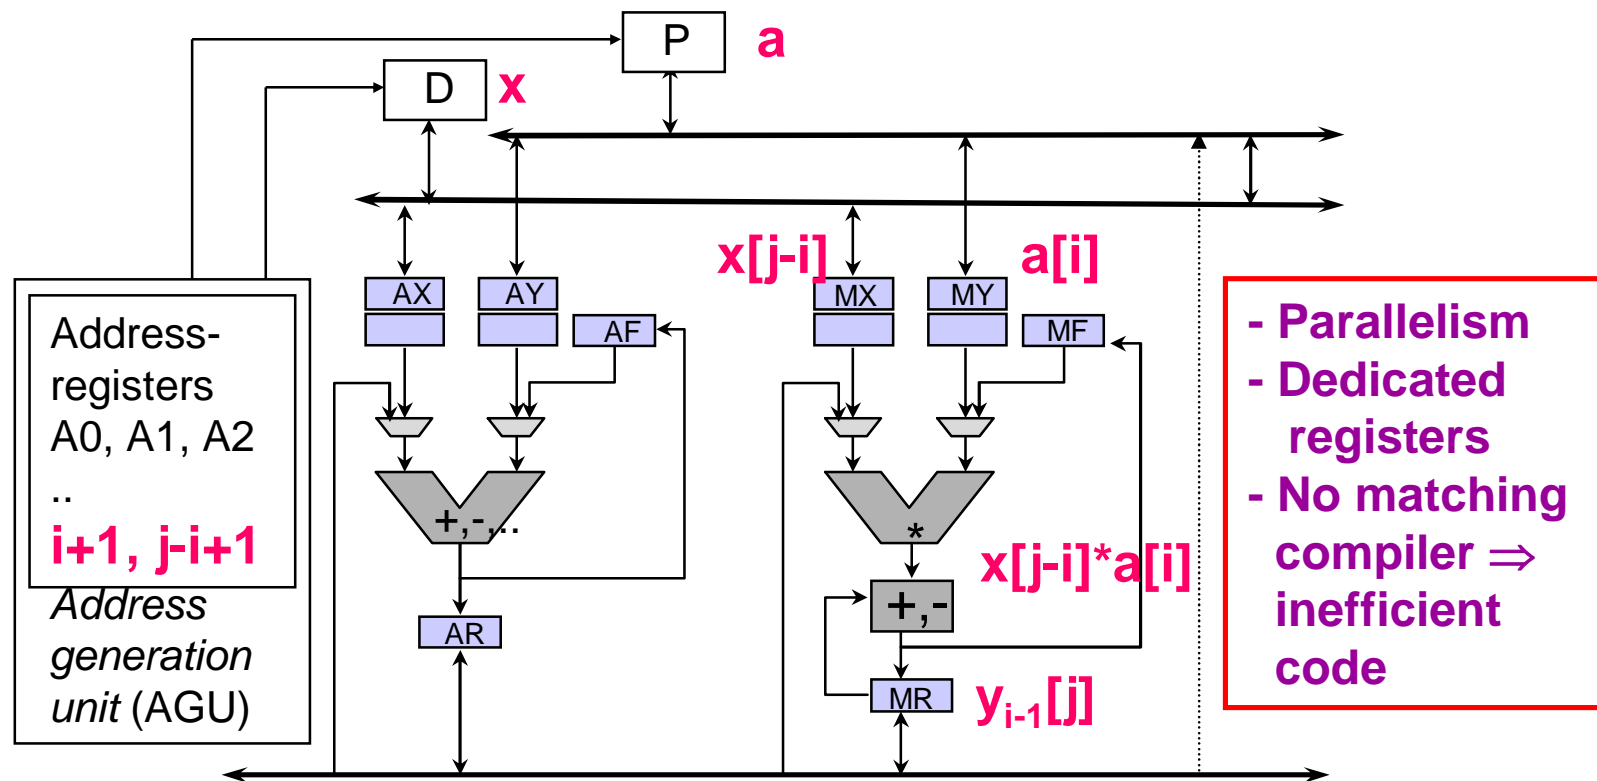


Reason for compiler-problems: Application-oriented Architectures

Application: u.a.: $y[j] = \sum_{i=0}^n x[j-i] * a[i]$

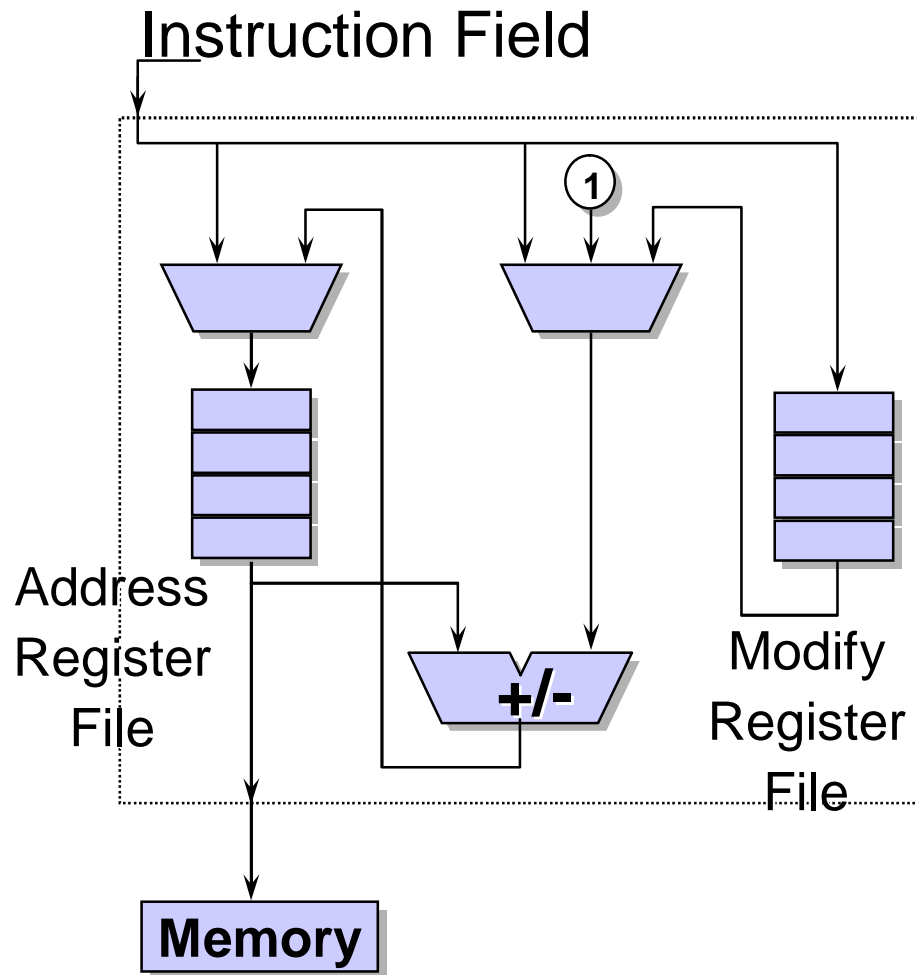
$$\forall i: 0 \leq i \leq n: y_i[j] = y_{i-1}[j] + x[j-i] * a[i]$$

Architecture: Example: Data path ADSP210x



Exploitation of parallel address computations

Generic address generation unit (AGU) model



Parameters:

k = # address registers

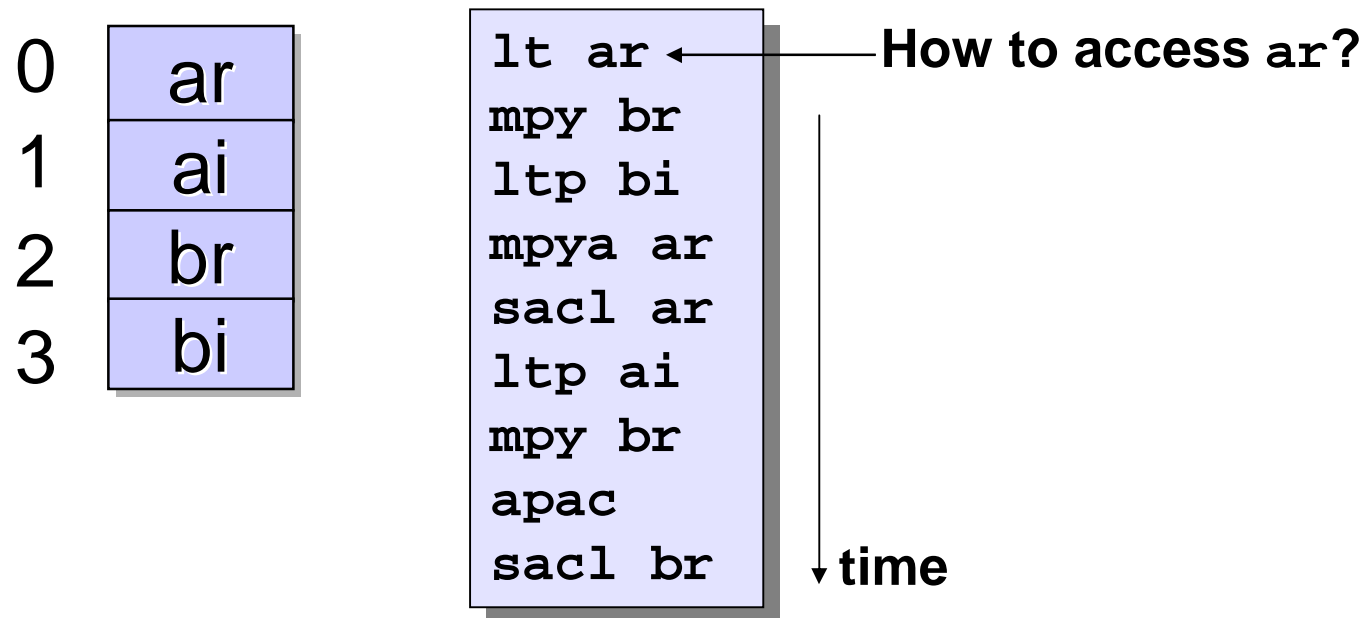
m = # modify registers

Cost metric for AGU operations:

Operation	cost
immediate AR load	1
immediate AR modify	1
auto-increment/ decrement	0
AR += MR	0

Address pointer assignment (APA)

Given: Memory layout + assembly code (without address code)



Address pointer assignment (APA) is the sub-problem of finding an allocation of address registers for a given memory layout and a given schedule.


General approach: Minimum Cost Circulation Problem

Let $G = (V, E, u, c)$, with (V, E) : directed graph

- $u: E \rightarrow \mathbb{R}_{\geq 0}$ is a capacity function,
- $c: E \rightarrow \mathbb{R}$ is a cost function; $n = |V|$, $m = |E|$.

Definition:

1. $g: E \rightarrow \mathbb{R}_{\geq 0}$ is called a **circulation** if it satisfies :
 $\forall v \in V: \sum_{w \in V: (v,w) \in E} g(v,w) = \sum_{w \in V: (w,v) \in E} g(w,v)$ (flow conservation)
2. g is **feasible** if $\forall (v,w) \in E: g(v,w) \leq u(v,w)$ (capacity constraints)
3. The cost of a circulation g is $c(g) = \sum_{(v,w) \in E} c(v,w) g(v,w)$.
4. There may be a lower bound on the flow through an edge.
5. The **minimum cost circulation problem** is to find a feasible circulation of minimum cost.




[K.D. Wayne: A Polynomial Combinatorial Algorithm for Generalized Minimum Cost Flow, <http://www.cs.princeton.edu/~wayne/papers/ratio.pdf>]

Mapping APA to the Minimum Cost Circulation Problem

Assembly sequence*

```
lt ar
mpy br
ltp bi
mpy ai
mpya ar
sac1 ar
ltp ai
mpy br
apac
sac1 br
```

* C2x pro-
cessor from ti



Flow into and out of variable nodes must be 1. Replace variable nodes by edge with lower bound=1 to obtain pure circulation problem

circulation selected


additional edges
of original graph
(only samples
shown)

[C. Gebotys: DSP Address
Optimization Using A Minimum
Cost Circulation Technique,
ICCAD, 1997]

Results according to Gebotys

Optimized code size

Original code size




Beyond basic blocks: - handling array references in loops -

Example:

```
for (i=2; i<=N; i++)  
{ .. B[i+1]    /*A2++ */  
  .. B[i]      /*A1-- */  
  .. B[i+2]    /*A2++ */  
  .. B[i-1]    /*A1++ */  
  .. B[i+3]    /*A2-- */  
  .. B[i] }    /*A1++ */
```

Cost for crossing loop boundaries considered.

Reference: A. Basu, R. Leupers, P. Marwedel: Array Index Allocation under Register Constraints, Int. Conf. on VLSI Design, Goa/India, 1999



Offset assignment problem (OA)

- Effect of optimised memory layout -

Let's assume that we can modify the memory layout

👉 offset assignment problem.

(k,m,r) -OA is the problem of generating a memory layout which minimizes the cost of addressing variables, with

👉 k : number of address registers

👉 m : number of modify registers

👉 r : the offset range

The case $(1,0,1)$ is called simple offset assignment (SOA),
the case $(k,0,1)$ is called general offset assignment (GOA).

SOA example

- Effect of optimised memory layout -

Variables in a basic block: Access sequence:

$V = \{a, b, c, d\}$

$S = (b, d, a, c, d, c)$

0	a	Load AR,1 ;b
1	b	AR += 2 ;d
2	c	AR -= 3 ;a
3	d	AR += 2 ;c
		AR ++ ;d
		AR -- ;c


cost: 4

0	b	Load AR,0 ;b
1	d	AR ++ ;d
2	c	AR +=2 ;a
3	a	AR -- ;c
		AR -- ;d
		AR ++ ;c


cost: 2

SOA example: Access sequence, access graph and Hamiltonian paths

access sequence: b d a c d c



access graph



maximum weighted path=
max. weighted Hamilton
path covering (MWHC)

0	b
1	d
2	c
3	a

memory layout

SOA used as a building block for more complex situations

➡ significant interest in good SOA algorithms



[Bartley, 1992; Liao, 1995]

Naïve SOA

Nodes are added in the order
in which they are used in the program.

Example:

Access sequence: $S = (b, d, a, c, d, c)$



0	b
1	d
2	a
3	c

memory layout

Liao's algorithm


Similar to Kruskal's spanning tree algorithms:

1. Sort edges of access graph $G=(V,E)$ according to their weight
2. Construct a new graph $G'=(V',E')$, starting with $E'=0$
3. Select an edge e of G of highest weight; If this edge does not cause a cycle in G' and does not cause any node in G' to have a degree > 2 then add this node to E' otherwise discard e .
4. Goto 3 as long as not all edges from G have been selected and as long as G' has less than the maximum number of edges ($|V|-1$).


Example: Access sequence: $S=(b, d, a, c, d, c)$

\nearrow \searrow \searrow \searrow \searrow \searrow
 1 0 1 0 0 0

Implicit edges of weight 0 for all unconnected nodes.





2 (c,d)
 1 (a,c)
 1 (a,d)
 1 (b,d)

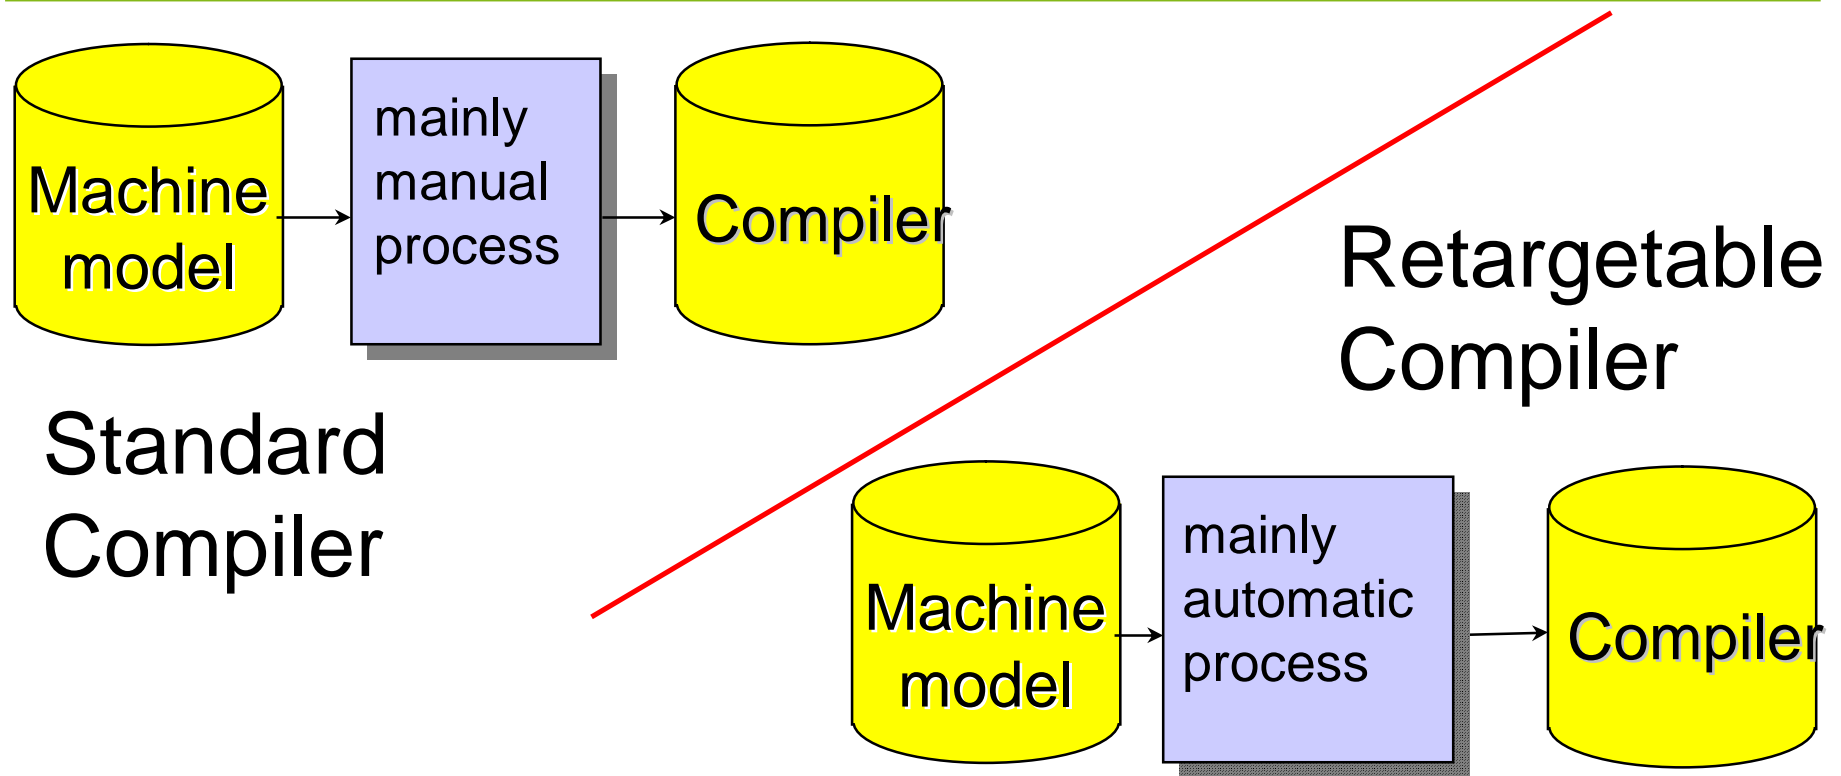


Liao's algorithm on a more complex graph

a b c d e f a d a d a c d f a d



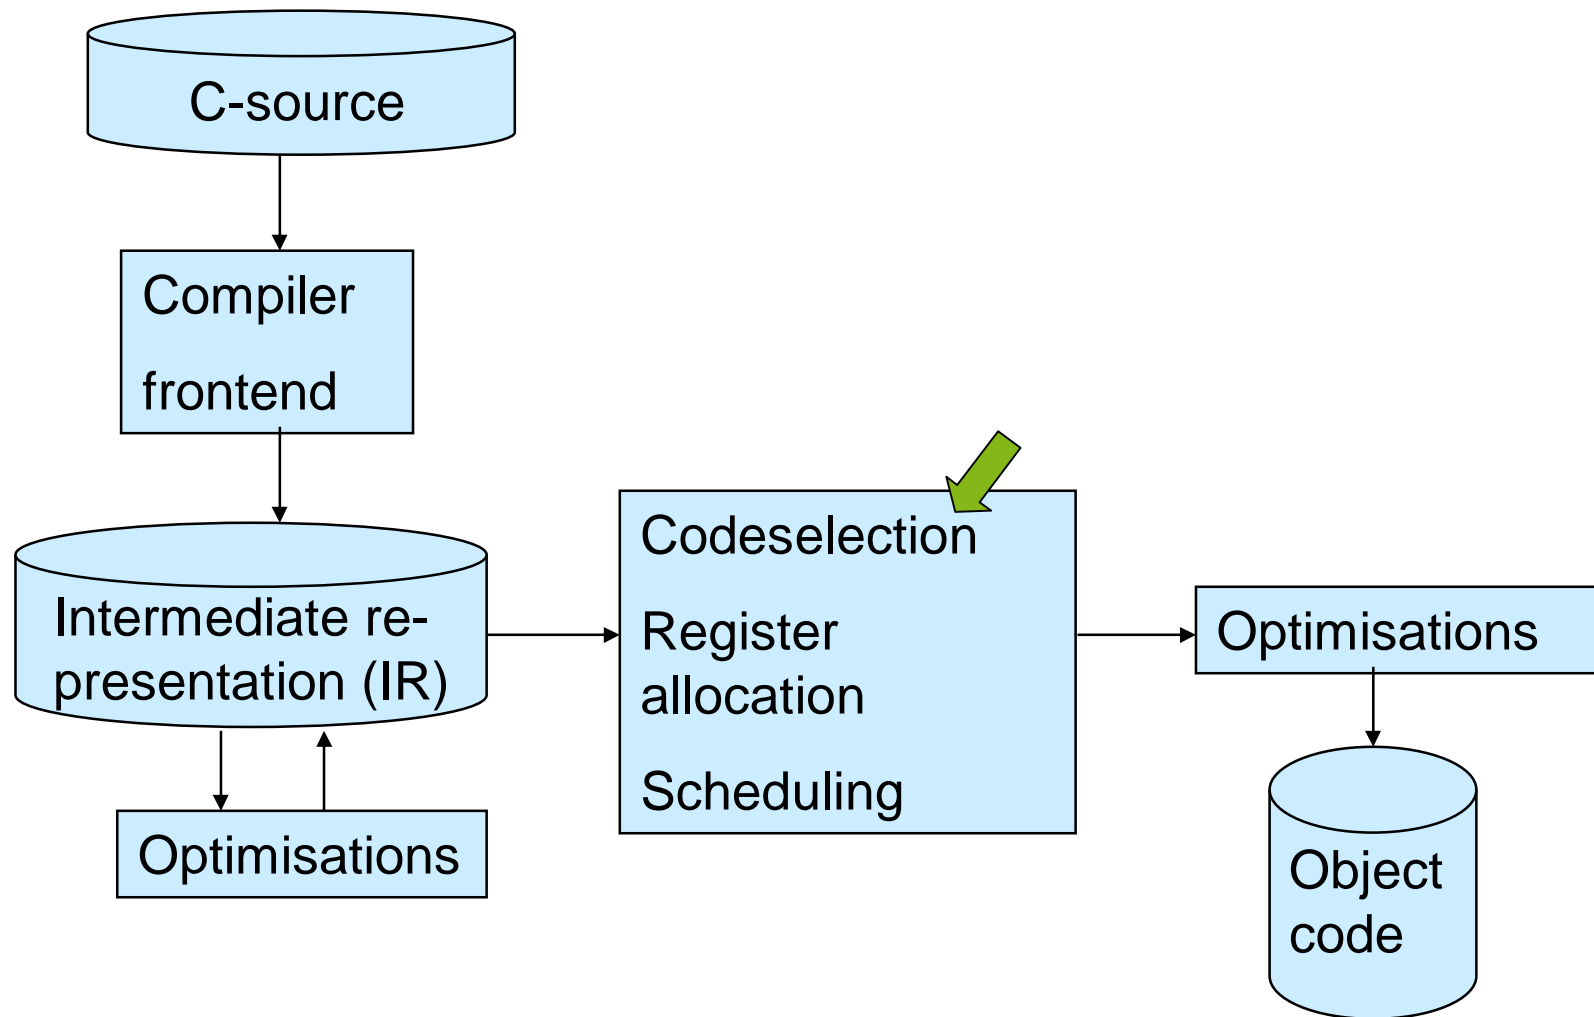
Retargetable Compilers vs. Standard Compilers




Developer retargetability: compiler specialists responsible for retargeting compilers.

User retargetability: users responsible for retargeting compiler.

Compiler structure



Code selection = covering DFGs



Does not yet
consider data
moves to input
registers.

Code selection by tree parsing (1)

Specification of grammar for generating a iburg parser*:

terminals: {MEM, *, +}

non-terminals: {reg1,reg2,reg3}

start symbol: reg1

rules:

“add” (cost=2): reg1 -> + (reg1, reg2)

“mul” (cost=2): reg1 -> * (reg1,reg2)

“mac” (cost=3): reg1 -> + (*(reg1,reg2), reg3)

“load” (cost=1): reg1 -> MEM

“mov2”(cost=1): reg2 -> reg1

“mov3”(cost=1): reg3 -> reg1

Code selection by tree parsing (2)

- nodes annotated with (register/pattern/cost)-triples -

“load”(cost=1):
reg1 -> MEM

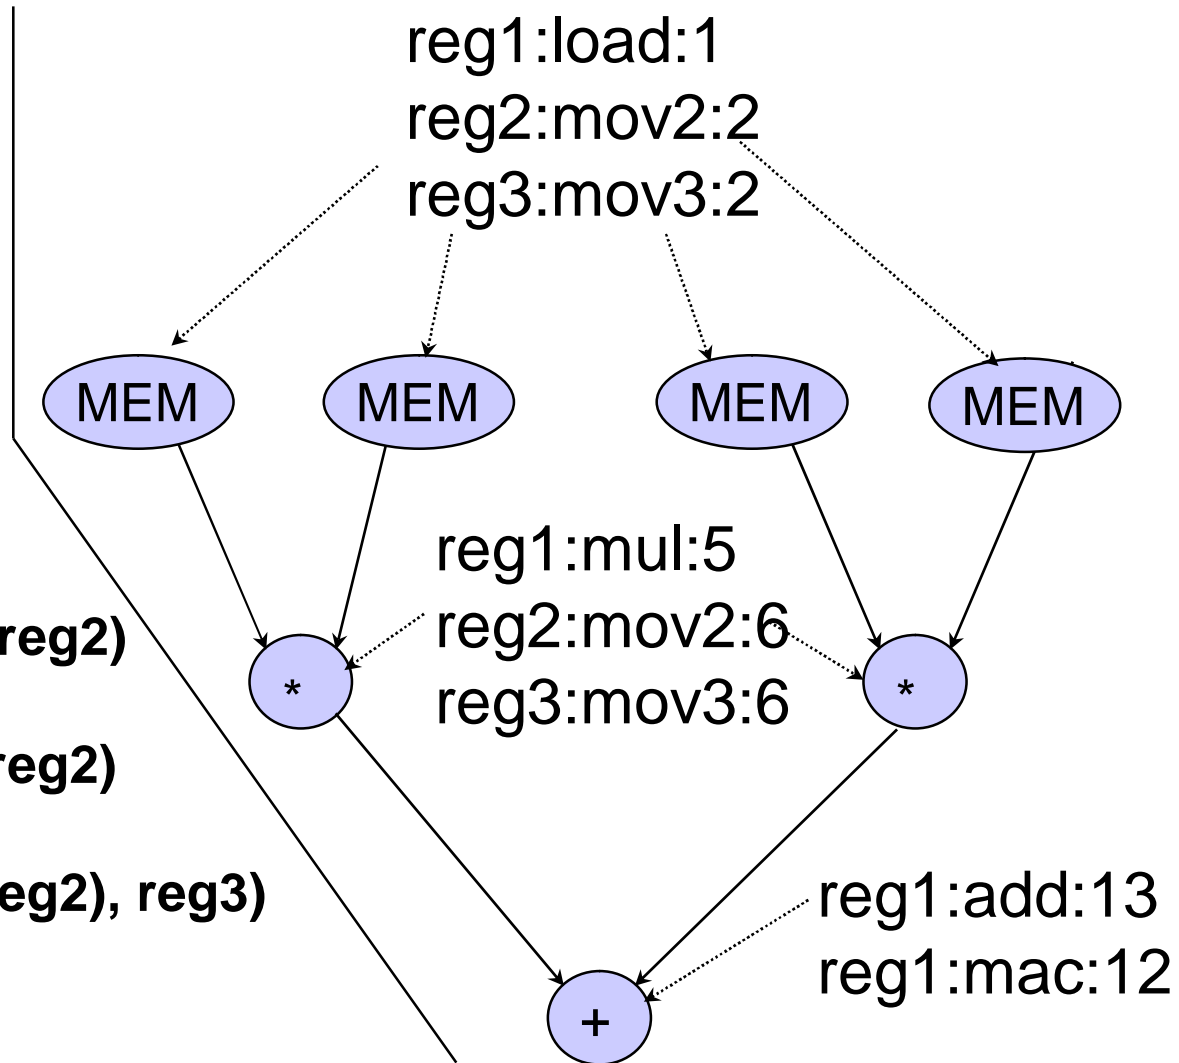
“mov2”(cost=1):
reg2 -> reg1

“mov3”(cost=1):
reg3 -> reg1

“add”(cost=2):
reg1 -> +(reg1, reg2)


“mul”(cost=2):
reg1 -> *(reg1, reg2)

“mac”(cost=3):
reg1 -> +(* (reg1, reg2), reg3)



Code selection by tree parsing (3)

- final selection of cheapest set of instructions -



Dynamic Voltage Scaling


Peter Marwedel
TU Dortmund
Informatik 12
Germany

2009/01/17



Voltage Scaling and Power Management

Dynamic Voltage Scaling



Recap from chapter 3: Fundamentals of dynamic voltage scaling (DVS)

Power consumption of CMOS circuits (ignoring leakage):

$$P = \alpha C_L V_{dd}^2 f \text{ with}$$

α : switching activity

C_L : load capacitance

V_{dd} : supply voltage

f : clock frequency

Delay for CMOS circuits:

$$\tau = k C_L \frac{V_{dd}}{(V_{dd} - V_t)^2} \text{ with}$$

V_t : threshold voltage

(V_t substantially < than V_{dd})


☞ Decreasing V_{dd} reduces P quadratically, while the run-time of algorithms is only linearly increased (ignoring the effects of the memory system).

Example: Processor with 3 voltages

Case a): Complete task ASAP

Task that needs to execute 10^9 cycles within 25 seconds.


V_{dd} [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
f_{max} [MHz]	50	40	25
cycle time [ns]	20	25	40



Case b): Two voltages

V_{dd} [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
f_{max} [MHz]	50	40	25
cycle time [ns]	20	25	40


b) $[V^2]$ 750M cycles @ 50 MHz + 250M cycles @ 25



$$\begin{aligned} E_b &= 750 \cdot 10^6 \times 40 \cdot 10^{-9} + \\ &\quad 250 \cdot 10^6 \times 10 \cdot 10^{-9} \text{ [J]} \\ &= 32.5 \text{ [J]} \end{aligned}$$

Case c): Optimal voltage

V_{dd} [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
f_{max} [MHz]	50	40	25
cycle time [ns]	20	25	40



Observations

- A minimum energy consumption is achieved for the ideal supply voltage of 4 Volts.
- In the following: **variable voltage processor** = processor that allows **any** supply voltage up to a certain maximum.
- It is expensive to support truly variable voltages, and therefore, actual processors support only a few fixed voltages.

Generalisation

Lemma [Ishihara, Yasuura]:

- If a variable voltage processor completes a task before the deadline, the energy consumption can be reduced.
- If a processor uses a single supply voltage V and completes a task T just at its deadline, then V is the unique supply voltage which minimizes the energy consumption of T .
- If a processor can only use a number of discrete voltage levels, then the two voltages which (almost*) minimize the energy consumption are the two immediate neighbors of the ideal voltage V_{ideal} possible for a variable voltage processor.

* Except for small amounts of energy resulting from the fact that cycle counts must be integers.

The case of multiple tasks: Assignment of optimum voltages to a set of tasks

N : the number of tasks

EC_j : the number of executed cycles of task j

L : the number of voltages of the target processor

V_i : the i^{th} voltage, with $1 \leq i \leq L$

F_i : the clock frequency for supply voltage V_i

T : the global deadline at which all tasks must have been completed

$X_{i,j}$: the number of clock cycles task j is executed at voltage V_i

SC_j : the average switching capacitance during the execution of task j (SC_j comprises the actual capacitance CL and the switching activity α)

Designing an ILP model

Simplifying assumptions of the ILP-model include the following:

- There is one target processor that can be operated at a limited number of discrete voltages.
- The time for voltage and frequency switches is negligible.
- The worst case number of cycles for each task are known.

Energy Minimization using an Integer Linear Programming Model


Minimize
$$E = \sum_{j=1}^N \sum_{i=1}^L SC_j \cdot x_{i,j} \cdot V_i^2$$

subject to
$$\forall j: \sum_{i=1}^L x_{i,j} = EC_j$$

and
$$\sum_{i=1}^L \sum_{j=1}^N \frac{x_{i,j}}{F_i} \leq T$$

Experimental Results

3 tasks; $EC_j=50 \cdot 10^9$; $SC_1=50$ pF; $SC_2=100$ pF; $SC_3=150$ pF





Task-level concurrency management


- The dynamic behavior of applications getting more attention.
- Energy consumption reduction is the main target.
- Some classes of applications (i.e. video processing) have a considerable variation in processing power requirements depending on input data.
- Static design-time methods becoming insufficient.
- Runtime-only methods not feasible for embedded systems.

→ How about mixed approaches?


Example of a mixed TCM




Static (compile-time) methods can ensure WCET feasible schedules, but waste energy in the average case.



...or they can define a probability for violating the deadline.



Mixed methods use compile-time analysis to define a set of possible execution parameters for each task.



Runtime scheduler selects the most energy saving, deadline preserving combination.

[IMEC, Belgium, <http://www.imec.be/>]

Dynamic power management (DPM)

Dynamic Power management tries to assign optimal power saving states.

- Questions: When to go to an power-saving state?
Different, but typically complex models:
- Markov chains, renewal theory ,

Summary

- Worst-case execution time aware compilation
- The offset assignment problem
- Retargetability and code selection
- Dynamic voltage scaling (DVS)
 - An ILP model for voltage assignment in a multi-tasking system
- Dynamic power management (DPM) (briefly)